

📖 Complete Code Breakdown - Line by Line

Table of Contents

1. [threat_detector.py - Main Model](#threat-detector)
2. [dashboard.py - Web Application](#dashboard)
3. [generate_sample_data.py - Test Data](#sample-data)
4. [dashboard.html - Web Interface](#html)

1 threat_detector.py - The Main Brain

Part 1: Imports (Lines 1-15)

```
```python
```

....

AI-Powered Threat Detection System with Explainability

....

...

\*\*What it does:\*\* This is a docstring - a description of what this file does. It doesn't execute, just documents.

```
```python
```

```
import os
```

...

What it does: Imports the `os` module for working with files and folders

Example: Creating folders, checking if files exist

```
```python
```

```
import numpy as np
```

...

\*\*What it does:\*\* Imports NumPy for math operations on arrays  
\*\*Think of it as:\*\* A super-powered calculator for lots of numbers at once  
\*\*Example:\*\* `np.mean([1,2,3])` calculates average = 2

```
```python
import pandas as pd
...```

```

What it does: Imports Pandas for working with tables of data (like Excel)
Think of it as: Excel in Python
Example: `pd.read_csv('file.csv')` reads a CSV file into a table

```
```python
import matplotlib.pyplot as plt
...
```


**What it does:** Imports Matplotlib for creating charts and graphs  
**Think of it as:** Making visual charts like bar graphs, line graphs  
**Example:** Creates the confusion matrix image


```

```
```python
import seaborn as sns
...
```


**What it does:** Imports Seaborn for prettier charts (built on top of Matplotlib)  
**Think of it as:** Matplotlib with automatic nice colors and styling


```

```
```python
from sklearn.model_selection import train_test_split
...
```


**What it does:** Imports function to split data into training and testing sets  
**Think of it as:** Dividing flashcards into "practice pile" and "test pile"  
**Example:** 70% for learning, 30% for testing


```

```
```python
from sklearn.preprocessing import StandardScaler, LabelEncoder
```
**What it does:** Imports tools to prepare data
- **StandardScaler:** Makes all numbers have similar scale (0-1 range)
- **LabelEncoder:** Converts text labels to numbers
**Example:** "DDoS" → 1, "BENIGN" → 0
```

```
```python
from sklearn.ensemble import RandomForestClassifier
```
**What it does:** Imports the Random Forest algorithm
**Think of it as:** A team of 100 decision trees voting together
**Analogy:** Like asking 100 experts and taking majority vote
```

```
```python
from sklearn.metrics import classification_report, confusion_matrix
```
**What it does:** Imports tools to measure how good the model is
**Think of it as:** Grading the model's test results
**Example:** Accuracy, precision, recall scores
```

```
```python
import shap
```
**What it does:** Imports SHAP for explaining predictions
**Think of it as:** The "why" explainer - shows which features mattered most
**Example:** "High packet rate (+0.3) made it predict DDoS"
```

```
```python
import joblib
```

---

\*\*What it does:\*\* Imports tool to save and load the trained model

\*\*Think of it as:\*\* Saving your trained model like saving a video game

\*\*Example:\*\* Save once, load later without retraining

```python

```
import warnings
```

```
warnings.filterwarnings('ignore')
```

What it does: Turns off warning messages to keep output clean

Think of it as: Muting non-critical notifications

```python

```
plt.switch_backend('Agg')
```

---

\*\*What it does:\*\* Tells Matplotlib to save images to files instead of showing on screen

\*\*Why:\*\* Works better in scripts that don't have a display (like servers)

---

## Part 2: The ThreatDetectionSystem Class (Lines 17-50)

```python

```
class ThreatDetectionSystem:
```

"""

Main class for network threat detection with explainability

"""

What it does: Creates a blueprint (class) for our threat detection system

Think of it as: A recipe that contains all methods for building the detector

Analogy: Like a car blueprint that shows all parts and how they work

```
```python
 def __init__(self, model_type='random_forest'):
 ...

 What it does: Initializes (sets up) the class when you create it

 Think of it as: The constructor - runs automatically when you create an instance

 Example: `detector = ThreatDetectionSystem()` ← This runs `__init__`
```

```
```python
    self.model_type = model_type
    self.model = None
    ...

    **What it does:**

    - `self.model_type` stores which ML algorithm to use (default: random_forest)
    - `self.model = None` creates empty placeholder for the trained model

    **Think of it as:** Setting up empty boxes to fill later
```

```
```python
 self.scaler = StandardScaler()
 self.label_encoder = LabelEncoder()
 ...

 What it does:

 - Creates a scaler object (normalizes feature values)
 - Creates label encoder (converts text to numbers)

 Example:

 - Scaler: 1000 → 0.5, 2000 → 1.0 (normalized)
 - Encoder: "DDoS" → 1, "BENIGN" → 0
```

```
```python
    self.feature_names = None
    self.explainer = None
```

What it does: Creates placeholders for:

- Feature names (list of column names)

- SHAP explainer object

Think of it as: Empty variables to fill during training

Part 3: Load and Preprocess Data Method (Lines 52-85)

```python

```
def load_and_preprocess_data(self, filepath):
```

"""

Load CICIDS2017 dataset and perform preprocessing

"""

...

\*\*What it does:\*\* Defines a method to load CSV file and clean it

\*\*Parameters:\*\* `filepath` = location of CSV file

\*\*Returns:\*\* Cleaned dataframe and label column name

```python

```
    print("[*] Loading dataset...")
```

```
    print(f"[*] Reading file: {filepath}")
```

...

What it does: Prints status messages so user knows what's happening

Think of it as: Progress updates

Output: `[*] Loading dataset...`

```python

```
if not os.path.exists(filepath):
```

```
 raise FileNotFoundError(f"Dataset not found at: {filepath}")
```

```

What it does: Checks if file exists; if not, shows error

Think of it as: Safety check before trying to open file

Example: If you typed wrong path, it tells you immediately

```python

```
df = pd.read_csv(filepath, encoding='utf-8')
```

```

What it does: Reads CSV file into a Pandas DataFrame (table)

Think of it as: Opening an Excel file

Result: `df` contains all rows and columns from CSV

Visual representation:

```

CSV File:      DataFrame (df):

| Port | Size | Label  | Port | Size | Label  |
|------|------|--------|------|------|--------|
| 80   | 500  | BENIGN | 80   | 500  | BENIGN |
| 443  | 600  | DDoS   | 443  | 600  | DDoS   |
| 80   | 550  | BENIGN | 80   | 550  | BENIGN |
| ...  |      |        |      |      |        |

```

```python

```
print(f"[+] Dataset shape: {df.shape}")
```

```

What it does: Prints how many rows and columns

Example output: `[+] Dataset shape: (225745, 79)` = 225,745 rows, 79 columns

```python

```
df = df.replace([np.inf, -np.inf], np.nan)
...
What it does: Replaces infinity values with NaN (Not a Number)
Why: Math errors can create infinity; we need to handle them
Example: Division by zero → infinity → NaN
```

```
```python
df = df.fillna(0)
...
**What it does:** Replaces all NaN (missing) values with 0
**Why:** Machine learning models can't handle missing data
**Example:** If "Packet Size" is missing → set to 0
```

```
```python
df = df.drop_duplicates()
...
What it does: Removes exact duplicate rows
Why: Duplicates can bias the model
Example: If same row appears 10 times, model thinks it's very important
```

```
```python
df.columns = df.columns.str.strip()
...
**What it does:** Removes extra spaces from column names
**Example:** `` Destination Port `` → ``Destination Port``
```

```
```python
label_col = 'Label' if 'Label' in df.columns else df.columns[-1]
...
What it does: Finds which column contains the labels (answers)
Logic:
```

- If column named "Label" exists, use it

- Otherwise, use the last column

\*\*Result:\*\* `label\_col = "Label"`

```python

```
print(f"[+] Label column: {label_col}")
```

```
print(f"[+] Attack types distribution:")
```

```
print(df[label_col].value_counts())
```

What it does: Shows how many of each attack type

Example output:

```
[+] Label column: Label
```

```
[+] Attack types distribution:
```

```
BENIGN 150000
```

```
DDoS 45000
```

```
PortScan 15000
```

```python

```
return df, label_col
```

---

\*\*What it does:\*\* Sends back the cleaned dataframe and label column name

\*\*Think of it as:\*\* Function's output that other code can use

---

## Part 4: Feature Engineering Method (Lines 87-110)

```python

```
def feature_engineering(self, df, label_col):
```

.....

Extract and engineer features from raw data

.....

...

What it does: Separates data into features (X) and labels (y)

Think of it as: Separating questions from answers

```python

```
print("\n[*] Performing feature engineering...")
```

...

\*\*What it does:\*\* Prints status update

\*\*The `\\n`:\*\* Creates a blank line before printing

```python

```
X = df.drop(columns=[label_col])
```

```
y = df[label_col]
```

...

What it does:

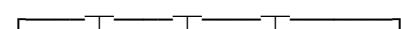
- 'X' = All columns EXCEPT the label (the features/inputs)

- 'y' = Only the label column (the answers/outputs)

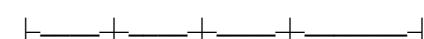
Visual:

...

Original DataFrame:



| Port | Size | IAT | Label |
|------|------|-----|-------|
|------|------|-----|-------|



| | | | |
|----|-----|-----|--------|
| 80 | 500 | 100 | BENIGN |
|----|-----|-----|--------|

→ X: Port, Size, IAT

| | | | |
|-----|-----|----|------|
| 443 | 600 | 50 | DDoS |
|-----|-----|----|------|

→ y: BENIGN, DDoS, BENIGN

| | | | |
|----|-----|-----|--------|
| 80 | 550 | 120 | BENIGN |
|----|-----|-----|--------|



```

```python

```
    numeric_cols = X.select_dtypes(include=[np.number]).columns
```

```
    X = X[numeric_cols]
```

```

\*\*What it does:\*\* Keeps only numeric columns (removes text columns)

\*\*Why:\*\* ML models need numbers, not text

\*\*Example:\*\* Keeps "Port: 80" but removes "Protocol: TCP"

```python

```
    self.feature_names = X.columns.tolist()
```

```

\*\*What it does:\*\* Saves column names in a list for later use

\*\*Example:\*\* `['Destination Port', 'Flow Duration', 'Total Fwd Packets', ...]`

```python

```
    y_encoded = self.label_encoder.fit_transform(y)
```

```

\*\*What it does:\*\* Converts text labels to numbers

\*\*Process:\*\*

1. `fit`: Learn the unique labels (BENIGN, DDoS, PortScan)
2. `transform`: Convert them to numbers

\*\*Example:\*\*

```

Before: After:

['BENIGN'] → [0]

['DDoS'] → [1]

['BENIGN'] → [0]

['PortScan'] → [2]

```python

```
print(f"[+] Features extracted: {len(self.feature_names)}")
```

```
print(f"[+] Classes: {list(self.label_encoder.classes_)}")
```

---

**\*\*What it does:\*\*** Shows how many features and what classes exist

**\*\*Example output:\*\***

---

```
[+] Features extracted: 77
```

```
[+] Classes: ['BENIGN', 'DDoS', 'PortScan']
```

---

```python

```
return X, y_encoded
```

****What it does:**** Returns the features (X) and encoded labels (y)

Part 5: Train Model Method (Lines 112-145)

```python

```
def train_model(self, X_train, y_train, n_estimators=100):
```

```
 """
```

Train the threat detection model

```
 """
```

---

**\*\*What it does:\*\*** Trains the Random Forest model on the training data

**\*\*Parameters:\*\***

- `X\_train`: Training features

- `y\_train`: Training labels
- `n\_estimators=100`: Number of decision trees (default 100)

```
```python
    print("\n[*] Training Random Forest model...")
    print(f"[*] Training samples: {len(X_train)}")
````
```

\*\*What it does:\*\* Shows training progress

\*\*Example:\*\* `[\*] Training samples: 157021`

```
```python
    X_train_scaled = self.scaler.fit_transform(X_train)
````
```

\*\*What it does:\*\* Normalizes the training data

\*\*Process:\*\*

1. `fit`: Learn the mean and std of each feature
2. `transform`: Scale values to standard range

\*\*Example:\*\*

...

Before scaling:      After scaling:

Packet Size: 1500 → 0.5

Packet Size: 3000 → 1.0

Packet Size: 750 → 0.25

Formula:  $(\text{value} - \text{mean}) / \text{std}$

...

\*\*Why scale?\*\*

- Feature "Bytes" might be 1,000,000
- Feature "Packets" might be 10

- Without scaling, model focuses only on large numbers
- After scaling, all features have equal importance

```python

```
self.model = RandomForestClassifier(
    n_estimators=n_estimators,
    max_depth=20,
    min_samples_split=10,
    random_state=42,
    n_jobs=-1,
    class_weight='balanced',
    verbose=0
)
```

What it does: Creates the Random Forest model with specific settings

Each parameter explained:

- `n_estimators=100`: Create 100 decision trees
 - More trees = more accurate but slower
 - Like asking 100 experts instead of 1
- `max_depth=20`: Each tree can be maximum 20 levels deep
 - Prevents overfitting (memorizing training data)
 - Like limiting how detailed questions can get
- `min_samples_split=10`: Need at least 10 samples to split a node
 - Prevents tiny splits on few examples
 - Ensures decisions are based on enough data
- `random_state=42`: Random seed for reproducibility
 - Using 42 always gives same random results

- Makes experiments repeatable
- `n_jobs=-1`: Use all CPU cores for parallel processing
 - Trains faster by using multiple processors
 - -1 means "use all available cores"
- `class_weight='balanced'`: Handle imbalanced data
 - If 90% BENIGN, 10% DDoS, model might always predict BENIGN
 - This balances importance of rare classes
- `verbose=0`: Don't print training progress details

****How Random Forest works:****

...

Tree 1: Port=80? → Size>500? → Predict: DDoS

Tree 2: Size>600? → IAT<100? → Predict: DDoS

Tree 3: Port=443? → Size>550? → Predict: BENIGN

...

Tree 100: (makes its prediction)

Final Prediction: Majority vote

- 70 trees say DDoS
 - 30 trees say BENIGN
- Prediction: DDoS (70%)

...

```python

```
self.model.fit(X_train_scaled, y_train)
```

```

****What it does:**** Actually trains the model (the learning happens here)

****Process:****

1. Model looks at features and labels
2. Finds patterns (e.g., high packet rate → DDoS)
3. Builds 100 decision trees
4. Each tree learns slightly different patterns

What's happening inside:

...

For each tree:

1. Randomly select features
2. Find best way to split data
3. Keep splitting until max_depth
4. Store the tree structure

...

Time: Takes 2-5 minutes depending on data size

```python

```
print("[+] Model training completed")
return X_train_scaled
```

...

\*\*What it does:\*\* Confirms training finished and returns scaled data

---

## Part 6: Evaluate Model Method (Lines 147-190)

```python

```
def evaluate_model(self, X_test, y_test):
```

"""

Evaluate model performance

"""

```

\*\*What it does:\*\* Tests how accurate the model is on unseen data

```python

```
print("\n[*] Evaluating model...")  
X_test_scaled = self.scaler.transform(X_test)
```

```

\*\*What it does:\*\* Scales test data using SAME scaling as training

\*\*Important:\*\* Use `transform` (not `fit\_transform`)

\*\*Why:\*\* We scale test data the same way as training data

\*\*Example:\*\*

```

Training: Mean=1000, Std=200

Test value: 1200

Scaled: $(1200 - 1000) / 200 = 1.0$

We DON'T recalculate mean/std for test data!

```

```python

```
y_pred = self.model.predict(X_test_scaled)  
y_pred_proba = self.model.predict_proba(X_test_scaled)
```

```

\*\*What it does:\*\*

- `y\_pred`: Predicted class (0, 1, 2)
- `y\_pred\_proba`: Probability for each class

\*\*Example:\*\*

```

Input: Network flow features

```
y_pred: 1 (DDoS)  
y_pred_proba: [0.05, 0.90, 0.05]  
BENIGN DDoS PortScan
```

Interpretation: 90% confident it's DDoS

...

```
```python  
print("\n" + "="*60)
print("CLASSIFICATION REPORT")
print("="*60)
print(classification_report(y_test, y_pred,
 target_names=self.label_encoder.classes_))
```

...

\*\*What it does:\*\* Prints detailed accuracy metrics

\*\*Example output:\*\*

...

```
=====
```

CLASSIFICATION REPORT

```
=====
```

precision recall f1-score support

BENIGN	0.99	0.98	0.99	45000
DDoS	0.97	0.99	0.98	13500
PortScan	0.96	0.95	0.96	4500

accuracy	0.98	63000
----------	------	-------

...

\*\*What each metric means:\*\*

- **Precision:** When model says "DDoS", how often is it right?
- Formula: True Positives / (True Positives + False Positives)
- Example: Predicted 100 as DDoS, 97 actually were = 97% precision
  
- **Recall:** Of all actual DDoS attacks, how many did we catch?
- Formula: True Positives / (True Positives + False Negatives)
- Example: 100 real DDoS, caught 99 = 99% recall
  
- **F1-Score:** Balance between precision and recall
- Formula:  $2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$
- Example: Good when both precision and recall are high
  
- **Support:** Number of actual examples in test set

```python

```
cm = confusion_matrix(y_test, y_pred)
```

```

**What it does:** Creates confusion matrix (table showing predictions vs reality)

**Visual explanation:**

...

Confusion Matrix:

Predicted

BENIGN DDoS PortScan

Actual BENIGN [44100] [600] [300] ← 44,100 correctly predicted

DDoS [ 200][13300] [ 0] ← 13,300 correctly predicted

PortScan[ 100][ 50] [4350] ← 4,350 correctly predicted

Diagonal = Correct predictions

Off-diagonal = Mistakes

...

```
```python
    plt.figure(figsize=(10, 8))

    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=self.label_encoder.classes_,
                yticklabels=self.label_encoder.classes_)

```

```

**\*\*What it does:\*\*** Creates visual heatmap of confusion matrix

**\*\*Parameters:\*\***

- `figsize=(10, 8)`: Size of image in inches
- `annot=True`: Show numbers in cells
- `fmt='d'`: Format numbers as integers (not decimals)
- `cmap='Blues'`: Color scheme (blue gradient)
- `xticklabels/yticklabels`: Labels for axes

```
```python
    plt.title('Threat Detection Confusion Matrix')

    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')

    plt.tight_layout()

```

```

**\*\*What it does:\*\*** Adds title and axis labels, adjusts spacing

```
```python
    output_path = 'outputs\\confusion_matrix.png'

    os.makedirs('outputs', exist_ok=True)

    plt.savefig(output_path, dpi=300, bbox_inches='tight')

    plt.close()

```

```

**\*\*What it does:\*\***

1. Sets output filename

2. Creates `outputs` folder if it doesn't exist
3. Saves image at 300 DPI (high quality)
4. Closes the plot to free memory

**\*\*Parameters explained:\*\***

- `dpi=300`: Dots per inch (higher = better quality)
- `bbox\_inches='tight'`: Remove extra whitespace
- `plt.close()`: Prevents memory leaks

```python

```
print(f"[+] Confusion matrix saved to '{output_path}'")
```

****What it does:**** Confirms file was saved

```python

```
self.plot_feature_importance()
return X_test_scaled, y_pred, y_pred_proba
```

---

**\*\*What it does:\*\***

1. Calls method to plot feature importance
2. Returns scaled test data and predictions

---

**## Part 7: Feature Importance Method (Lines 192-215)**

```python

```
def plot_feature_importance(self, top_n=20):
```

.....

Visualize feature importance

.....

```

\*\*What it does:\*\* Shows which features matter most for predictions

```python

```
if hasattr(self.model, 'feature_importances_'):  
    ...
```

What it does: Checks if model has feature importance scores

Why: Only tree-based models have this attribute

```python

```
importances = self.model.feature_importances_
indices = np.argsort(importances)[-top_n:]
...
```

\*\*What it does:\*\*

- Gets importance score for each feature
- Finds top 20 most important features
- `np.argsort`: Returns indices that would sort the array
- `[-top\_n:]`: Takes last 20 (highest values)

\*\*Example:\*\*

```

Features: ['Port', 'Size', 'Duration', 'IAT']

Importances: [0.05, 0.30, 0.15, 0.40]

Sorted indices: [0, 2, 1, 3]

Top 2: [1, 3] (Size and IAT)

```

```python

```
plt.figure(figsize=(10, 8))  
plt.barh(range(len(indices)), importances[indices], color='steelblue')  
...
```

****What it does:**** Creates horizontal bar chart

- `barh`: Horizontal bars (easier to read long names)
- `range(len(indices))`: Y-axis positions (0, 1, 2, ...)
- `importances[indices]`: X-axis values (importance scores)

```python

```
plt.yticks(range(len(indices)),
 [self.feature_names[i] for i in indices])
...
````
```

****What it does:**** Labels Y-axis with feature names

****List comprehension breakdown:****

```python

```
[self.feature_names[i] for i in indices]
```

Step by step:

```
for i in indices: # For each index
 self.feature_names[i] # Get feature name at that index
 # Add to list
```

Result: ['Feature 1', 'Feature 2', ...]

...

```python

```
plt.xlabel('Feature Importance')  
plt.title(f'Top {top_n} Most Important Features')  
plt.tight_layout()  
...  
````
```

**\*\*What it does:\*\*** Adds labels and adjusts layout

```python

```
output_path = 'outputs\\feature_importance.png'
```

```
plt.savefig(output_path, dpi=300, bbox_inches='tight')
plt.close()
print(f"[+] Feature importance saved to '{output_path}'")
````
```

\*\*What it does:\*\* Saves chart to file and confirms

---

## Part 8: SHAP Explainability (Lines 217-280)

```
```python
```

```
def initialize_explainer(self, X_train_scaled):
    """
    Initialize SHAP explainer for model interpretability
    """
````
```

\*\*What it does:\*\* Sets up SHAP to explain model predictions

```
```python
```

```
    print("\n[*] Initializing SHAP explainer...")
    self.explainer = shap.TreeExplainer(self.model)
````
```

\*\*What it does:\*\* Creates SHAP explainer for Random Forest

\*\*TreeExplainer:\*\* Optimized for tree-based models (fast and accurate)

\*\*How SHAP works:\*\*

---

For each prediction:

1. Calculate baseline (average prediction)
2. For each feature:
  - See how prediction changes with/without it

- Calculate contribution (SHAP value)
3. Sum all contributions = Final prediction

Example:

Baseline: 50% threat

+ Packet rate (+30%)

+ Port 80 (+10%)

- Normal size (-5%)

= 85% threat (DDoS)

...

```python

```
print("[+] SHAP explainer initialized")
```

```
return self.explainer
```

...

****What it does:**** Confirms setup and returns explainer

```python

```
def explain_prediction(self, sample_index, X_test_scaled, y_test, y_pred):
```

...

**\*\*What it does:\*\*** Explains why model made a specific prediction

**\*\*Parameters:\*\***

- `sample\_index`: Which test sample to explain

- `X\_test\_scaled`: Scaled test features

- `y\_test`: True labels

- `y\_pred`: Predicted labels

```python

```
print(f"\n[*] Explaining prediction for sample {sample_index}...")
```

```
sample = X_test_scaled[sample_index:sample_index+1]
```

...

****What it does:****

- Prints status
- Gets specific sample (row) from test data
- `'[sample_index:sample_index+1]` : Slicing to keep 2D shape

****Why [index:index+1]?****

```python

X[0] → 1D array [1.2, 0.5, 3.1]

X[0:1] → 2D array [[1.2, 0.5, 3.1]]

Model needs 2D input (even for 1 sample)

...

```python

```
shap_values = self.explainer.shap_values(sample)
```

...

****What it does:**** Calculates SHAP values for this sample

****Output:**** Array showing contribution of each feature

****Example:****

...

Features: ['Port', 'Size', 'Duration', 'IAT']

Values: [80, 1500, 120000, 50]

SHAP values: [+0.1, +0.3, -0.05, +0.4]

Interpretation:

- Port contributes +0.1 to threat probability
- Size contributes +0.3 (important!)
- Duration reduces threat by -0.05
- IAT contributes +0.4 (most important!)
- ...

```
```python
 true_label = self.label_encoder.classes_[y_test[sample_index]]
 pred_label = self.label_encoder.classes_[y_pred[sample_index]]
```
**What it does:** Converts numeric predictions back to text labels

**Example:**  

```
y_test[0] = 1 → true_label = "DDoS"
y_pred[0] = 1 → pred_label = "DDoS"
```

```

```
```python
print(f"\n{'='*60}")
print("PREDICTION EXPLANATION")
print(f"{'='*60}")
print(f"True Label: {true_label}")
print(f"Predicted Label: {pred_label}")
print(f"Prediction Correct: {true_label == pred_label}")
```
**What it does:** Prints comparison of true vs predicted label

**Example output:**  

=====
PREDICTION EXPLANATION
=====
True Label: DDoS
Predicted Label: DDoS
Prediction Correct: True
```

```

```
```python
if isinstance(shap_values, list) and len(shap_values) > 0:
    shap_vals = shap_values[y_pred[sample_index]]
else:
    shap_vals = shap_values
```

```

\*\*What it does:\*\* Handles multi-class SHAP values

\*\*Explanation:\*\*

- For binary: SHAP returns single array
- For multi-class: SHAP returns list of arrays (one per class)
- We select array for the predicted class

```
```python
```

```
try:
    plt.figure(figsize=(10, 6))
    expected_val = (self.explainer.expected_value[y_pred[sample_index]])
    if isinstance(self.explainer.expected_value, np.ndarray)
        else self.explainer.expected_value
```

```

\*\*What it does:\*\*

- Creates new figure
- Gets baseline/expected value (average prediction)
- Handles both array and single value cases

```
```python
```

```
shap.waterfall_plot(
    shap.Explanation(
        values=shap_vals[0],
        base_values=expected_val,
        data=sample[0],
        feature_names=self.feature_names
```

```

```
),
 show=False
)
...

```

\*\*What it does:\*\* Creates SHAP waterfall plot

\*\*Components:\*\*

- `values`: SHAP contributions for each feature
- `base\_values`: Starting point (average)
- `data`: Actual feature values
- `feature\_names`: Names for labels
- `show=False`: Don't display, just prepare to save

\*\*Waterfall plot visualization:\*\*

...

$E[f(X)] = 0.5$  (baseline)

$\downarrow +0.3$  (Packet rate high)  
 $\downarrow +0.1$  (Port 80)  
 $\downarrow -0.05$  (Normal duration)  
 $\downarrow +0.2$  (Low IAT)

$f(x) = 1.05 \rightarrow \text{DDoS}$  (105% probability, capped at 100%)

...

```python

```
plt.tight_layout()  
output_path = f'outputs\\explanation_sample_{sample_index}.png'  
plt.savefig(output_path, dpi=300, bbox_inches='tight')  
plt.close()  
print(f"[+] Explanation saved to '{output_path}'")  
...  
...
```

What it does: Saves the waterfall plot

```
```python
 except Exception as e:
 print(f"[!] Could not generate waterfall plot: {e}")
```
**What it does:** Catches and reports any errors
**Why:** SHAP plotting can sometimes fail; we don't want to crash
```

```
```python
 return shap_vals
```
**What it does:** Returns SHAP values for further analysis
```

```
## Part 9: Save/Load Model (Lines 282-320)
```

```
```python
<function_calls> <invoke name="artifacts"> <parameter name="command">update</parameter>
<parameter name="id">code_explanation_guide</parameter> <parameter
name="old_str">```python </parameter> <parameter name="new_str">def save_model(self,
filepath='threat_detector.pkl'): """ Save trained model and preprocessing objects """
```
**What it does:** Saves everything needed to use the model later
```

```
**Parameter:** `filepath` = where to save (default: threat_detector.pkl)
```

```
```python
 model_data = {
 'model': self.model,
 'scaler': self.scaler,
 'label_encoder': self.label_encoder,
 'feature_names': self.feature_names
 }
```

```

What it does: Creates dictionary with all components

Why save all these?

- `model`: The trained Random Forest
- `scaler`: To scale new data the same way
- `label_encoder`: To convert predictions back to text
- `feature_names`: To know feature order

Think of it as: Packing everything into a suitcase

```python

```
joblib.dump(model_data, filepath)
```

```

What it does: Serializes (saves) the dictionary to disk

File format: .pkl (pickle) - compressed binary format

Size: Usually 10-50 MB depending on model

What's inside the .pkl file:

threat_detector.pkl

```
|—— Random Forest (100 trees)  
|—— StandardScaler (mean/std for each feature)  
|—— LabelEncoder (class mappings)  
└—— Feature names (list of strings)
```

```python

```
print(f"[+] Model saved to '{filepath}'")
```

```

What it does: Confirms save completed

```python

```
def load_model(self, filepath='threat_detector.pkl'):
```

.....

## Load trained model

11

三

**\*\*What it does:\*\*** Loads previously saved model

```
```python
```

```
if not os.path.exists(filepath):  
    raise FileNotFoundError(f"Model file not found: {filepath}")
```

...

****What it does:**** Safety check - make sure file exists before trying to load

```
```python
```

```
model_data = joblib.load(filepath)
```

111

**\*\*What it does:\*\*** Deserializes (loads) the dictionary from disk.

**\*\*Time:\*\*** Fast, usually < 1 second

```
```python
```

```
self.model = model_data['model']
self.scaler = model_data['scaler']
self.label_encoder = model_data['label_encoder']
self.feature_names = model_data['feature_names']
```

三

****What it does:**** Unpacks dictionary and assigns to class attributes

****Think of it as:** Unpacking the suitcase**

```
```python
```

```
print(f"[+] Model loaded from '{filepath}'")
```

三

**\*\*What it does:\*\*** Confirms load completed

#### **\*\*Usage example:\*\***

```
```python
```

```
# Save after training
```

```
detector.save_model('my_model.pkl')
```

```
# Later, in a different script:  
  
detector = ThreatDetectionSystem()  
  
detector.load_model('my_model.pkl')  
  
# Now ready to make predictions without retraining!  
  
'''
```

- - -

Part 10: Main Execution Function (Lines 322-400)

```python

```
def main():
```

11

## Main execution pipeline

11

11

**\*\*What it does:\*\*** Orchestrates the entire training process

**\*\*Think of it as:\*\*** The conductor of an orchestra

```python

```
print("=*60)
```

```
print("AI-POWERED THREAT DETECTION SYSTEM")
```

```
print("Windows Edition")
```

```
print("="*60)
```

11

****What it does:**** Prints nice header banner

```
```python
```

```
detector = ThreatDetectionSystem(model_type='random_forest')
```

111

**\*\*What it does:\*\*** Creates instance of our class

**\*\*Equivalent to:\*\***

```
```python
```

```

detector = ThreatDetectionSystem() # Using default
```
What happens: Runs `__init__` method, sets up empty model
```python
filepath = r'data\raw\Friday-WorkingHours-Afternoon-DDos.pcap_ISCX.csv'
```
What it does: Sets path to dataset
The `r` before string: Raw string (treats backslashes literally)
Why: In Windows, paths use backslashes
```python
try:
    df, label_col = detector.load_and_preprocess_data(filepath)
except FileNotFoundError:
    print("\n[!] Dataset not found!")
    print("[!] Please:")
    print("  1. Download CICIDS2017")
    print("  2. Place CSV file in: data\\raw\\")
    print("  3. Or run: python src\\generate_sample_data.py")
    return
```
What it does: Tries to load data; if fails, shows helpful error
try/except: Error handling
- `try`: Attempt this code
- `except`: If error occurs, do this instead
- `return`: Exit function early

```

**Flow:**

Try to load data



Success? → Continue



Failure? → Show error message → Exit

```
```python
if len(df) > 100000:
    print(f"\n[*] Large dataset detected. Sampling 100,000 rows...")
    df = df.sample(n=100000, random_state=42)
````
```

\*\*What it does:\*\* If dataset is huge, use smaller sample

\*\*Why:\*\* Training on 2.8M rows takes hours; 100K is faster for testing

\*\*`random\_state=42`:\*\* Ensures same sample every time

```
```python
```

```
X, y = detector.feature_engineering(df, label_col)
```

```
````
```

\*\*What it does:\*\* Separates features and labels

\*\*Returns:\*\*

- `X`: DataFrame with features
- `y`: Array with encoded labels

```
```python
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)
```

```
````
```

\*\*What it does:\*\* Splits data into training (70%) and testing (30%) sets

\*\*Parameters:\*\*

- `test\_size=0.3`: 30% for testing
- `random\_state=42`: Reproducible split
- `stratify=y`: Maintain class distribution in both sets

\*\*Stratify explanation:\*\*

Original data:

70% BENIGN, 30% DDoS

Without stratify:

Train: 80% BENIGN, 20% DDoS ← Unbalanced!

Test: 50% BENIGN, 50% DDoS

With stratify:

Train: 70% BENIGN, 30% DDoS ← Balanced!

Test: 70% BENIGN, 30% DDoS

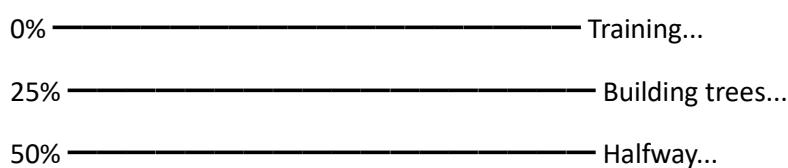
```
```python
print(f"\n[+] Training set: {X_train.shape}")
print(f"[+] Test set: {X_test.shape}")
```
What it does: Shows split sizes
Example: `'[+] Training set: (70000, 77)`
```

```
```python
X_train_scaled = detector.train_model(X_train, y_train, n_estimators=100)
```
What it does: Trains the model
```

\*\*What's happening:\*\*

1. Scales training data
2. Creates Random Forest with 100 trees
3. Fits model to training data
4. Returns scaled data

\*\*Timeline:\*\*



75% ━━━━━━━━━━━━━━━━━━━━━━━ Almost done...

100% ━━━━━━━━━━━━ Complete!

```python

```
X_test_scaled, y_pred, y_pred_proba = detector.evaluate_model(X_test, y_test)
```

...

What it does: Tests model accuracy

Returns:

- `X_test_scaled`: Scaled test features
- `y_pred`: Predicted classes
- `y_pred_proba`: Prediction probabilities

```python

```
detector.initialize_explainer(X_train_scaled)
```

...

\*\*What it does:\*\* Sets up SHAP for explanations

```python

```
detector.explain_prediction(0, X_test_scaled, y_test, y_pred)
```

```
detector.explain_prediction(10, X_test_scaled, y_test, y_pred)
```

...

What it does: Explains 2 example predictions (samples 0 and 10)

Why 2? Shows variety - one might be correct, one incorrect

```python

```
detector.generate_shap_summary(X_test_scaled)
```

...

\*\*What it does:\*\* Creates global SHAP summary plot

```python

```
detector.save_model('threat_detector.pkl')
```

...

What it does: Saves trained model to disk

Result: File `threat_detector.pkl` created in root folder

```
```python
print("\n" + "="*60)
print("ANALYSIS COMPLETE")
print("="*60)
print("\nGenerated files:")
print(" - threat_detector.pkl (trained model)")
print(" - outputs\\confusion_matrix.png")
print(" - outputs\\feature_importance.png")
print(" - outputs\\explanation_sample_*.png")
print(" - outputs\\shap_summary.png")
````
```

****What it does:**** Shows summary of what was created

```
```python
print("\nNext step: Run the dashboard")
print(" python src\\dashboard.py")
````
```

****What it does:**** Tells user what to do next

```
```python
if __name__ == "__main__":
 main()
````
```

****What it does:**** Runs main() only when script is executed directly

****Why needed:**** Allows importing functions without running main()

****Explanation:****

```
```python
When you run: python threat_detector.py
__name__ == "__main__" → True → Runs main()

When you import: from threat_detector import ThreatDetectionSystem
__name__ == "threat_detector" → False → Doesn't run main()
````
```

This completes the line-by-line explanation of `threat_detector.py`!

****Summary of what this file does:****

1. Loads and cleans network traffic data
2. Trains Random Forest classifier
3. Evaluates model performance
4. Generates SHAP explanations
5. Creates visualizations (charts/plots)
6. Saves trained model to disk

****Next file:**** `dashboard.py` - The web application

```</parameter>